

Education

University of Utah

PHD IN COMPUTER SCIENCE AND ENGINEERING

- Advised by Prof. Saday Sadayappan
- Research interest: High performance computing for sparse tensor contractions

Salt Lake City, USA

Aug. 2021 - Present

Birla Institute of Technology and Science, Pilani

BACHELOR OF ENGINEERING

Major: Computer Science

Pilani, India

Aug. 2015 - Dec. 2018

Publications

- [1] V. T. Chakaravarthy, S. S. Pandian, **Raje, Saurabh**, and Y. Sabharwal. On optimizing distributed non-negative tucker decomposition. In *Proceedings of the ACM International Conference on Supercomputing (ICS)*, pages 238–249, 2019.
- [2] V. T. Chakaravarthy, S. S. Pandian, **Raje, Saurabh**, Y. Sabharwal, T. Suzumura, and S. Ubaru. Efficient scaling of dynamic graph neural networks. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis, SC '21*, New York, NY, USA, 2021. Association for Computing Machinery.
- [3] S. Goyal, A. R. Choudhury, **Raje, Saurabh**, V. Chakaravarthy, Y. Sabharwal, and A. Verma. PoWER-BERT: Accelerating BERT inference via progressive word-vector elimination. In H. D. III and A. Singh, editors, *Proceedings of the 37th International Conference on Machine Learning (ICML)*, volume 119 of *Proceedings of Machine Learning Research*, pages 3690–3699, Virtual, 13–18 Jul 2020. PMLR.
- [4] S. Islam, S. Balasubramaniam, P. Goyal, A. Sultana, L. Bhutani, **Raje, Saurabh**, and N. Goyal. A rapid prototyping approach for high performance density-based clustering. In *2019 IEEE International Conference on Data Science and Advanced Analytics (DSAA)*, pages 260–269. IEEE, 2019.
- [5] A. Kannan, A. Roy Choudhury, V. Saxena, **Raje, Saurabh**, P. Ram, A. Verma, and Y. Sabharwal. Hyperaspo: Fusion of model and hyper parameter optimization for multi-objective machine learning. In *2021 IEEE International Conference on Big Data (Big Data)*, pages 790–800, 2021.
- [6] S. E. Kurt, **Raje, Saurabh**, A. Sukumaran-Rajam, and P. Sadayappan. Sparsity-aware tensor decomposition. In *2022 IEEE International Parallel and Distributed Processing Symposium (IPDPS)*, pages 952–962, 2022.
- [7] **Raje, Saurabh**, A. Goel, S. Sharma, K. Aggarwal, D. Mantri, and T. Kumar. Development of on board computer for a nanosatellite. *68th International Astronautical Congress (IAC)*, 2017.
- [8] **Raje, Saurabh**, S. Vaderia, N. Wilson, and R. Panigrahi. Decentralised firewall for malware detection. In *2017 International Conference on Advances in Computing, Communication and Control (ICAC3)*, pages 1–5. IEEE, 2017.
- [9] **Saurabh Raje** and F. Wagner. Kvik: A task based middleware with composable scheduling policies, 2020.
- [10] Y. Xu, **Raje, Saurabh**, A. Rountev, G. Sabin, A. Sukumaran-Rajam, and P. Sadayappan. Training of deep learning pipelines on memory-constrained gpus via segmented fused-tiled execution. In *Proceedings of the 31st ACM SIGPLAN International Conference on Compiler Construction, CC 2022*, page 104–116, New York, NY, USA, 2022. Association for Computing Machinery.

Work Experience

University of Utah

DOCTORAL RESEARCHER

ACCELERATING SPARSE LINEAR ALGEBRA

Salt Lake City, Utah

August 2019 - Present

- Currently working on new representations for sparse tensors with domain specific patterns.
- In collaboration with Pacific Northwest National labs, this research aims to accelerate quantum chemistry simulations.
- Co-developed a novel implementation for sparse-tensor decomposition (**SpTL**).
- **SpTL** reduces data movement and load imbalance to beat the state-of-the-art run-time.
- Co-developed a system to train convolutional neural networks (CNN)s on large images (20,000 x 20,000).
- This effectively tiles the dataflow through CNNs to enable processing of massive images on a single GPU system.

IBM Research

RESEARCH ENGINEER

ACCELERATING AI

Delhi, India

August 2019 - August 2021

- Worked with the model compression team to make AI faster.
- Designed *PowerBERT*, a new model that is up to **4.5x faster** than **BERT** for inference.
- This work was published in **ICML'20**, and was integrated into IBM OneNLP product stack.
- Implemented a new method to train massive Graph Neural Networks faster using supercomputers (published at **SC'21**)
- Implemented novel representations for sparse tensors. This was used to accelerate the tucker decomposition algorithm.
- Co-invented 4 **patents** on model compression techniques and multiobjective optimisation.

ETH Zurich

SCIENTIFIC ASSISTANT

COMPILERS FOR DEEP LEARNING

Zurich, Switzerland

March 2019 - August 2019

- Accelerated the training of Deep Neural Networks using the **DACE** language developed in-house.
- DACE is a domain specific language for HPC workloads that uses a novel Stateful Dataflow Graph (SDFG) based Intermediate Representation.
- Wrote a Tensorflow frontend for DACE that parses a TF computation graph to build a DACE SDFG.
- Added a pattern based compiler transformation on the IR to reduce GPU kernel calls and repetitive memory access.
- Achieved at-par performance for ResNet-50 in comparison to Tensorflow and CuDNN.

INRIA

BACHELOR THESIS

MIDDLEWARE FOR PARALLEL PROGRAMMING

Grenoble, France

September 2018 - February 2019

- Developed **Kvik**: a task based middleware in the **Rust** language.
- **Kvik** makes sequential code run in parallel without significant changes, by creating independent tasks.
- In particular, it provides tunable task splitting strategies that can be composed with each other.
- Wrote the fastest parallel merge sort using **Kvik** (2.5x faster than Intel TBB for 50 threads).

BITS Pilani

RESEARCH ASSISTANT

PARALLELIZING COMPILERS

Pilani, India

August 2017 - December 2018

- Worked on the DWARF domain specific language compiler.
- The compiler generates parallel code with MPI calls for various data mining applications.
- Modelled the data dependencies for density based and hierarchical clustering algorithms.
- Built a new optimisation layer that increased the granularity of parallelism.

Honors & Awards

- 2021 **Winner**, Patent Plateau Award - IBM India Research Lab
- 2021 **Winner**, Outstanding Technical Achievement Award - IBM India Research Lab
- 2020 **Winner**, Distinguished Paper Award - IBM India Research Lab
- 2020 **Winner**, Awesome Team Award - IBM India Research Lab
- 2018 **Winner**, Best Poster Award - IBM India Research Lab
- 2017 **Winner**, Mercedes Benz Hack.Banglore 2018
- 2016 **Winner**, Best Paper Award - APOGEE (BITS Pilani's technical festival)

Presentations

Mobile World Congress 2018

PRESENTER FOR DAIMLER AG

Barcelona, Spain

February 2018

- Invited by Daimler AG to present our winning hackathon prototype.
- The prototype was built to detect pedestrians using low cost IR sensors.
- This would allow for level 4+ automated driving.

Skills

Languages	Rust, Python, C, C++, Java
Frameworks	PyTorch, Tensorflow, Caffe, CuDNN, Git
HPC Libraries	openMPI, openMP, Intel TBB