Saurabh Raje

saurabh.raje@utah.edu | saurabhmraje.com | smr97 | Saurabh Raje | saurabhmraje

Education

University of Utah
Salt Lake City, Utah

PhD in Computer Science and Engineering

Aug. 2021 - Present

Advised by Prof. Saday Sadayappan
- Thesis: Accelerating sparse tensor contractions

Birla Institute of Technology and Science, Pilani
Pilani, India

Bachelor of Engineering

Aug. 2015 - Dec. 2018

Major: Computer Science
- Thesis: Kvik: A task based middleware with composable scheduling policies
- In collaboration with DATAMOVE team at INRIA Grenoble.

Research Statement

I am a final year PhD student doing research in High Performance Computing (HPC) and compilers. My focus is on optimizing sparse tensor contractions. Part of my work has been on hand-writing fast kernels for CPUs and GPUs to do tensor times tensor contractions. I also developed compiler transformations to optimize chains of sparse tensor contractions, and have worked on scaling up sparse tensor decomposition algorithms.

Some of my previous research at IBM has been on accelerating training and inference of deep neural networks - specifically in NLP and graph-based learning. Other than hand-tuning kernels, I have also worked on domain specific languages (DSL) to generate fast code for CPUs and GPUs.

Work Experience

Apple
Cupertino, California

PhD Intern

BNNSGraph Compiler

May 2025 - August 2025

- Improved the BNNSGraph compiler that supports AI inference on CPUs.
- Added several IR transformations to fuse kernels and reduce copies
- This made some text classification models more than twice as fast.
- Lead to significant improvements on the GeekbenchAI benchmark for iPhones.

Apple
Cupertino, California

PhD Intern

BNNSGraph Compiler

May 2024 - August 2024

- Wrote a fused attention kernel for Apple CPUs.
- Added pattern based transformations to the BNNSGraph Compiler to use the kernel.
- Reduced data movement and made self attention significantly faster.
- The kernel is currently used for on device inference for several language models running on Apple devices.

University of Utah
Salt Lake City, Utah

Doctoral Researcher

Accelerating Sparse Linear Algebra

August 2019 - Present

- Wrote compiler optimizations and hand tuned kernels for sparse tensor contractions.
- Built CoNST - a code generator that fuses sparse contraction loops across multiple expressions.
- Wrote FaSTCC - a library of the fastest CPU kernels for sparse tensor times tensor contractions.
- Added a sparse backend to TAMM - a compiler used by computational chemists at Pacific Northwest National Labs
- Wrote GPU kernels for sparse tensor times tensor operations.
- Co-developed a PyTorch extension to train convolutional neural networks (CNN)s on large images (20,000 x 20,000).

IBM Research
Delhi, India

Research Engineer

Model compression and optimization

August 2019 - August 2021

- Designed **PowerBERT** - up to 4.5x faster than BERT for inference (published in **ICML'20**), and integrated into IBM OneNLP product stack.
- Implemented a new method to train massive Graph Neural Networks faster using supercomputers (published at **SC'21**)
- Implemented novel representations for sparse tensors. This was used to accelerate the tucker decomposition algorithm.
- Co-invented 4 patents on model compression techniques and multiobjective optimisation.

ETH Zurich
Zurich, Switzerland

Scientific Assistant

Compilers for Deep Learning

March 2019 - August 2019

- Accelerated the training of Deep Neural Networks using the DACE language developed in-house.
- DACE is a domain specific language for HPC workloads that uses a novel Stateful Dataflow Graph (SDFG) based Intermediate Representation.
- Wrote a Tensorflow frontend for DACE that parses a TF computation graph to build a DACE SDFG.
- Added a pattern based compiler transformation on the IR to reduce GPU kernel calls and repetitive memory access.
- Achieved at-par performance for ResNet-50 in comparison to Tensorflow and CuDNN.

INRIA
Grenoble, France

Bachelor Thesis

Middleware for Parallel Programming

September 2018 - February 2019

- Developed Kvik: a task based middleware in the Rust language.
- Kvik makes sequential code run in parallel without significant changes, by creating independent tasks.
- In particular, it provides tunable task splitting strategies that can be composed with each other.
- Wrote the fastest parallel merge sort using Kvik (2.5x faster than Intel TBB for 50 threads).

Selected Publications

Full list here

[1] **Raje, Saurabh**, H. McCoy, A. Rountev, P. Pandey, and P. Sadayappan. Fastcc: Fast sparse tensor contractions on cpus. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, To Appear **SC** '25, New York, NY, USA, 2025. Association for Computing Machinery.

[2] **Raje, Saurabh**, Y. Xu, A. Rountev, E. F. Valeev, and P. Sadayappan. Const: Code generator for sparse tensor networks. *ACM Trans. Archit. Code Optim.*, 21(4), Nov. 2024.

[3] S. Goyal, A. R. Choudhury, **Raje, Saurabh**, V. Chakaravarthy, Y. Sabharwal, and A. Verma. PoWER-BERT: Accelerating BERT inference via progressive word-vector elimination. In H. D. III and A. Singh, editors, *Proceedings of the 37th International Conference on Machine Learning (**ICML**)*, volume 119 of *Proceedings of Machine Learning Research*, pages 3690–3699, Virtual, 13–18 Jul 2020. PMLR.

[4] **Raje, Saurabh Manish**, S. Goyal, A. R. Choudhury, Y. Sabharwal, and A. Verma. Accelerating inference of neural network models via dynamic early exits, Nov. 10 2022. *US Patent* App. 17/307,501.

[5] V. T. Chakaravarthy, S. S. Pandian, **Raje, Saurabh**, Y. Sabharwal, T. Suzumura, and S. Ubaru. Efficient scaling of dynamic graph neural networks. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, **SC** '21, New York, NY, USA, 2021. Association for Computing Machinery.

[6] S. E. Kurt, **Raje, Saurabh**, A. Sukumaran-Rajam, and P. Sadayappan. Sparsity-aware tensor decomposition. In *2022 IEEE International Parallel and Distributed Processing Symposium (**IPDPS**)*, pages 952–962, 2022.

[7] V. T. Chakaravarthy, S. S. Pandian, **Raje, Saurabh**, and Y. Sabharwal. On optimizing distributed non-negative tucker decomposition. In *Proceedings of the ACM International Conference on Supercomputing (**ICS**)*, pages 238–249, 2019.

[8] Y. Xu, **Raje, Saurabh**, A. Rountev, G. Sabin, A. Sukumaran-Rajam, and P. Sadayappan. Training of deep learning pipelines on memory-constrained gpus via segmented fused-tiled execution. In *Proceedings of the 31st ACM SIGPLAN International Conference on Compiler Construction*, **CC** 2022, page 104–116, New York, NY, USA, 2022. Association for Computing Machinery.